

00 A CONCEPT, EXPLAINED —

The Mental Model.

The brain for company management and agentic AI.

Robert Blust · Software Engineer & Architect

RESTS ON `concept Mental Model`

A company's knowledge lives everywhere — and nowhere.

Vision, strategy, processes, roles, KPIs, rules — scattered across heads, documents and tools.

Neither people nor AI can rely on it. What if all of it were a graph a machine could read?

RESTS ON `concept Master`

One structured knowledge base — the **brain** of the company.

I call it the Mental Model: the facts as Markdown, structured by a meta-model.

Not new documentation but a machine-readable **projection of the company's operating system** — versioned like code, synced from the real systems.

RESTS ON

concept Mental Model

vision One model, true everywhere

Everything that makes a company — in one place.

Roles & Teams

who does what

KPIs

what we measure

Concepts

what we define

Processes & Gates

how work flows

Rules

what we enforce

Architecture decisions

what we decided

Strategy & Objectives

where we invest

Features

what we build


Values

what we believe

Versioned like code, machine-readable, **ONE source of truth.**

RESTS ON **concept Company of one**

Everything cross-references — no dead links.



Filename IDs

the canonical name of every entry

Exact-name references


no markdown links that rot

Validated on every change

a contradiction surfaces immediately

RESTS ON `concept Master` `rule The model is corrected first`

Governance a machine can reason over.



Owner

Head of Product · gate sign-off

Gate criterion

Tech-Debt Impact Score (-3 ... +3)

Rule

score > +2 → cleanup enters scope

Kill criterion

adoption < X% → deprecate

RESTS ON `decision Agents write the model, and I approve every change to it`

`rule The Owner's word merges`

Rules, metrics and decisions — structured and owned.

Rule

“Deployments follow dev → test → prod.” — Owner: Engineering.

KPI

“Lead Time for Changes” — lower is better. — Owner: Engineering.

Architecture decision

“JSONB only for contextual data.” — rejected: querying on it.

Every rule, metric and decision: **named, justified, owned.**

RESTS ON `rule` A check counts only when its output was seen `kpi` Change Lead Time `decision` An agent's commit is authored by the seat it held

The same model — for people and agents.

FOR PEOPLE

Onboarding by role, decisions, alignment.

“Who owns the deployment?” — answered from one source.

One source — **two audiences, no drift.**

RESTS ON

`surface blust.ch website`

`surface mcp.blust.ch MCP server`

FOR AGENTS

A context layer plus MCP tools.

Agents act inside vision, strategy, process and rules.

The right context makes AI far more reliable.

Without context, AI knows neither your decisions nor your rules — it guesses.
The Mental Model shrinks that space and keeps agents inside the guardrails.
Always with human oversight.

RESTS ON

kpi Grounded Answer Rate

surface chat.blust.ch chat

Explicit knowledge pays off twice — for people and for agents.

Structure is the lever. The cleaner the model, the better people lead
and the more reliably agents act.

RESTS ON `vision One model, true everywhere`